

In [129]:

```
# Import relevant libraries
import pandas as pd
import numpy as np
import statsmodels.api as sm
# Time Series Analysis Library: AR, VAR, ARMA; ACF&PACF
import statsmodels.tsa as tsa
# Import acf & pacf plotting functions
from statsmodels.graphics.tsaplots import plot_acf
from statsmodels.graphics.tsaplots import plot_pacf
# Plotting Library
import matplotlib
import matplotlib.pyplot as plt
# Use the lzip to pretty-print short descriptions for the test returns
from statsmodels.compat import lzip
# Ljung-Box test function
from statsmodels.stats.diagnostic import acorr_ljungbox
```

In [130]:

```
pip install arch
```

Requirement already satisfied: arch in c:\users\35387\anaconda3\lib\site-packages (4.15)
Requirement already satisfied: pandas>=0.23 in c:\users\35387\anaconda3\lib\site-packages (from arch) (1.0.5)
Note: you may need to restart the kernel to use updated packages.
Requirement already satisfied: statsmodels>=0.9 in c:\users\35387\anaconda3\lib\site-packages (from arch) (0.11.1)
Requirement already satisfied: numpy>=1.14 in c:\users\35387\anaconda3\lib\site-packages (from arch) (1.18.5)
Requirement already satisfied: cython>=0.29.14 in c:\users\35387\anaconda3\lib\site-packages (from arch) (0.29.21)
Requirement already satisfied: scipy>=1.0.1 in c:\users\35387\anaconda3\lib\site-packages (from arch) (1.5.0)
Requirement already satisfied: property-cached>=1.6.3 in c:\users\35387\anaconda3\lib\site-packages (from arch) (1.6.4)
Requirement already satisfied: pytz>=2017.2 in c:\users\35387\anaconda3\lib\site-packages (from pandas>=0.23->arch) (2020.1)
Requirement already satisfied: python-dateutil>=2.6.1 in c:\users\35387\anaconda3\lib\site-packages (from pandas>=0.23->arch) (2.8.1)
Requirement already satisfied: patsy>=0.5 in c:\users\35387\anaconda3\lib\site-packages (from statsmodels>=0.9->arch) (0.5.1)
Requirement already satisfied: six>=1.5 in c:\users\35387\anaconda3\lib\site-packages (from python-dateutil>=2.6.1->pandas>=0.23->arch) (1.15.0)

In [131]:

```
# Check auxiliary regression for DF test
from arch.unitroot import ADF
```

In [132]:

```
data = pd.read_csv('GDPC1.csv')
```

In [133]:

```
# 1. Rename the columns
data.rename(columns={'DATE': 'Date', 'GDPC1': 'Real GDP'}, inplace=True)
```

In [134]:

```
# 2. Setting the Date column of strings to the datetime format in Pandas:
data['Date'] = pd.to_datetime(data['Date'], format='%Y/%m/%d')
```

In [135]:

```
# 3. Set the Date column as Index of the dataframe
data.set_index('Date', inplace=True)
```

In [136]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 103 entries, 1995-01-01 to 2020-07-01
Data columns (total 1 columns):
 #   Column   Non-Null Count   Dtype  
--- 
 0   Real GDP    103 non-null    float64
dtypes: float64(1)
memory usage: 1.6 KB
```

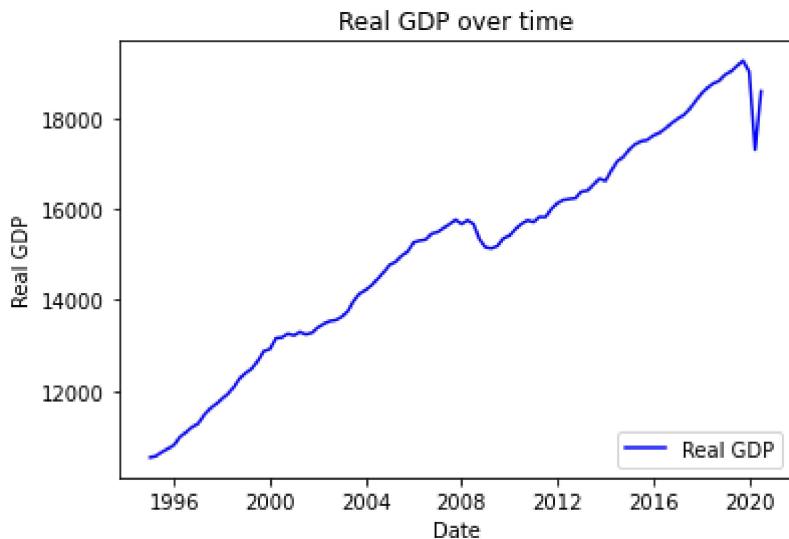
In [137]:

```
#1.1 Graphical Detection of Stationarity
```

In [138]:

```
#Plotting the original series against the time to detect stationarity
```

```
# Plotting the CPI series against the time index of dataframe
plt.plot(data, 'b-',label='Real GDP')
# Set the labels for axies
plt.xlabel('Date')
plt.ylabel('Real GDP')
# Set the title for graph
plt.title('Real GDP over time')
# Set the legend position
plt.legend(loc='lower right')
# Save the figure
plt.savefig('Real GDP Plot.jpg',bbox_inches='tight')
# Dispaly the figure
plt.show()
```



In [139]:

```
#Correlogram Plotting
```

In [140]:

```
#Applying the plot_acf() and plot_pacf()functions
```

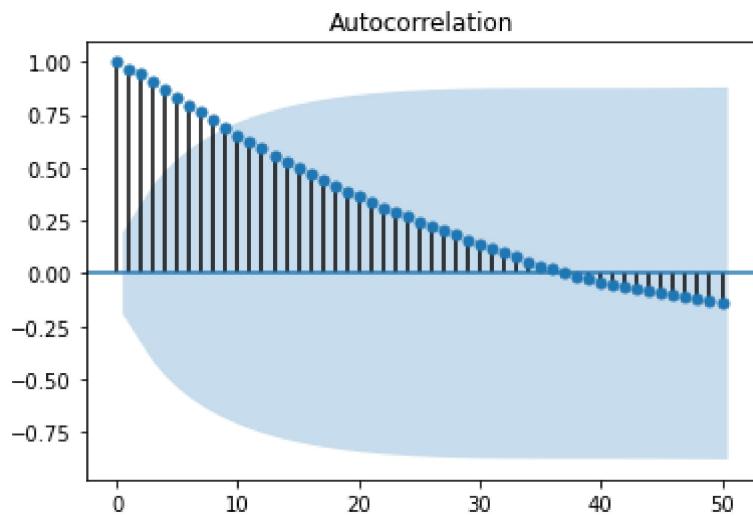
```
#Check whether the series is stationary or non-stationary:
```

#Stationary: Absence of correlation

#Non-stationary: Gradually declining correlogram (spurious correlation due to common tr

In [227]:

```
plot1 = plot_acf(data, lags=50)
plt.show()
```

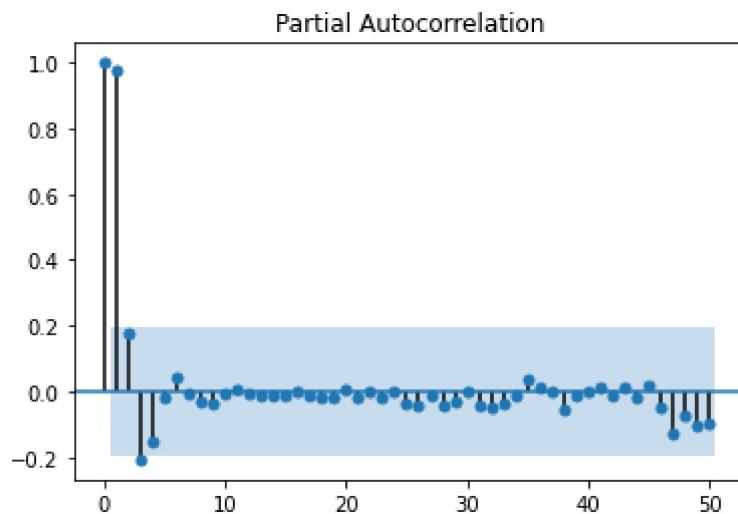


In [142]:

```
#We know this series non stationary as the AFC is slowly decreasing, implying a spurious co
#the empirical correlogram is useless in identifying the know stationary process.
```

In [226]:

```
# Plot the Pacf (CPI)
plot2 = plot_pacf(data, lags=50)
```



In [144]:

#1.2 Unit Roots Testing: Augmented Dickey-Fuller Test

#There are 3 verisons, however, all have the same null hypothesis that the series has a unit root

In [145]:

#1st Dickey-Fuller test for a random walk against a stationary autoregressive process of order 1

#H0 = Series has a unit root

#H1 = Series is stationary AR(1)

#Test stats not following the usual t-distribution of under the null: should compare it with standard normal distribution

In [146]:

1. Dickey-Fuller Test without a constant and trend

```
adf_1_name = ['Dickey-Fuller Test Statistic', 'p-value', 'Number of lags used', 'Number of observations',
              'Critical values', 'maximized information criterion']
```

Set regression parameter as 'nc': no constant, no trend

Number of lags is chosen to minimize the corresponding information criterion(AIC by default)

```
adf_1 = sm.tsa.stattools.adfuller(data['Real GDP'], regression='nc')
```

```
lzip(adf_1_name, adf_1)
```

#nc = alternative hypothesis

Out[146]:

```
[('Dickey-Fuller Test Statistic', 3.9359168647015883),
 ('p-value', 0.9999998246268814),
 ('Number of lags used', 1),
 ('Number of observations', 101),
 ('Critical values',
  {'1%': -2.5882321870404863,
   '5%': -1.9439589708250309,
   '10%': -1.614431731329}),
 ('maximized information criterion', 1230.296276131842)]
```

In [147]:

#p value really high so we cannot reject the null hypothesis

#thus we can say that this time series has a unit root and non stationary

In [148]:

```
#2. Dickey-Fuller test for a random walk against a stationary AR(1) with drift

# c = H1:

# 2. Dickey-Fuller Test with a constant
adf_2_name = ['Dickey-Fuller Test Statistic', 'p-value', 'Number of lags used', 'Number of ob
               'Critical values', 'maximized information criterion']
# Set regression parameter as 'c': constant only (default)
# Number of lags is chosen to minimize the corresponding information criterion(AIC by defau
adf_2 = sm.tsa.stattools.adfuller(data['Real GDP'], regression='c')
lzip(adf_2_name, adf_2)
```

Out[148]:

```
[('Dickey-Fuller Test Statistic', -1.7828088500834927),
 ('p-value', 0.38902940220094234),
 ('Number of lags used', 1),
 ('Number of observations', 101),
 ('Critical values',
  {'1%': -3.4968181663902103,
   '5%': -2.8906107514600103,
   '10%': -2.5822770483285953}),
 ('maximized information criterion', 1228.1991276243475)]
```

In [149]:

```
#Again the results show that we can reject the Null Hypothesis but not with as much certa
#Series has a unit root.
```

In [150]:

```
# 3. Dickey-Fuller Test with a constant and trend

adf_3_name = ['Dickey-Fuller Test Statistic', 'p-value', 'Number of lags used', 'Number of ob
               'Critical values', 'maximized information criterion']
# Set regression parameter as 'ct': constant and trend
adf_3 = sm.tsa.stattools.adfuller(data['Real GDP'], regression='ct')
lzip(adf_3_name, adf_3)
```

Out[150]:

```
[('Dickey-Fuller Test Statistic', -1.9395885077225477),
 ('p-value', 0.6339163353741536),
 ('Number of lags used', 1),
 ('Number of observations', 101),
 ('Critical values',
  {'1%': -4.051321648595896,
   '5%': -3.454889141998309,
   '10%': -3.1530564880069027}),
 ('maximized information criterion', 1226.0854437315197)]
```

In [151]:

```
#Again the results show that we can reject the Null Hypothesis but not with as much certa
#Series has a unit root.

#Conclusion: We have a non-stationary time series.

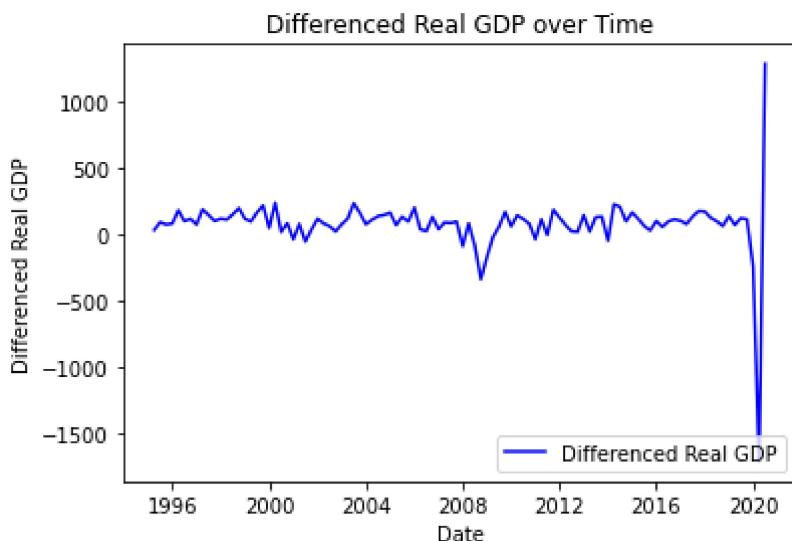
#Now we convert it to a stationary time series in order to fit an ARMA model to it.
```

In [152]:

```
# Take the first difference
diff_data = data.diff()
# Drop NA
diff_data.dropna(inplace=True)
```

In [153]:

```
# Plotting the CPI series against the time index of dataframe
plt.plot(diff_data, 'b-', label='Differenced Real GDP')
# Set the Labels for axies
plt.xlabel('Date')
plt.ylabel('Differenced Real GDP')
# Set the title for graph
plt.title('Differenced Real GDP over Time')
# Set the Legend position
plt.legend(loc='lower right')
# Save the figure
plt.savefig('Differenced Real GDP Plot.jpg', bbox_inches='tight')
# Display the figure
plt.show()
```



In [154]:

```
adf_dt_name = ['Dickey-Fuller Test Statistic', 'p-value', 'Number of lags used', 'Number of o
                 'Critical values', 'maximized information criterion']
# DF test for diff_data: differenced CPI
adf_dt = sm.tsa.stattools.adfuller(diff_data['Real GDP'])
lzip(adf_dt_name, adf_dt)
```

Out[154]:

```
[('Dickey-Fuller Test Statistic', -11.878973406830053),
 ('p-value', 6.24292235256828e-22),
 ('Number of lags used', 0),
 ('Number of observations', 101),
 ('Critical values',
  {'1%': -3.4968181663902103,
   '5%': -2.8906107514600103,
   '10%': -2.5822770483285953}),
 ('maximized information criterion', 1215.7726620333765)]
```

In [155]:

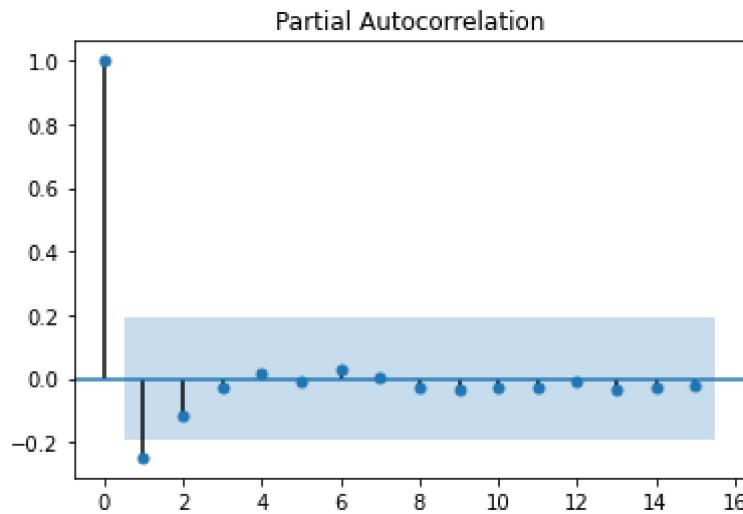
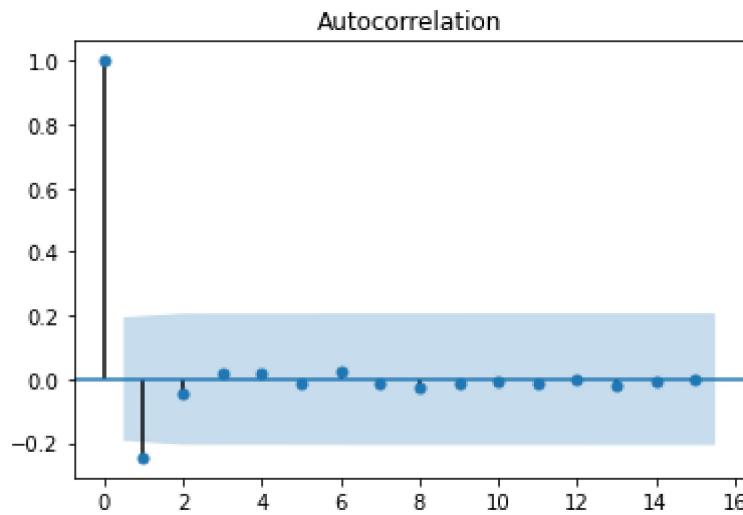
```
# DF Test statistic is negative and high and p value is basically zero so we reject the H0
# Time series is does not have a unit rate and therefore is stationary.
```

In [156]:

```
#ARMA MODEL ESTIMATION
```

In [193]:

```
plot_acf(diff_data, lags=15)
plot_pacf(diff_data, lags=15)
plt.show()
```



In [194]:

#There is one lag that is significantly different from zero. So, this would define the order in this case if would be of order 1.

#There is 1 relevant lag at which the PACF is significantly different from zero. The 1st La for the purpose of completing the process of comparing ARMA models.

#Thus we have 3 possible models to compare

```
#ARMA(1,1)
#ARMA(2,1)
#ARMA(7,1)
```

In [160]:

```
#ARMA (1,1)
```

In [189]:

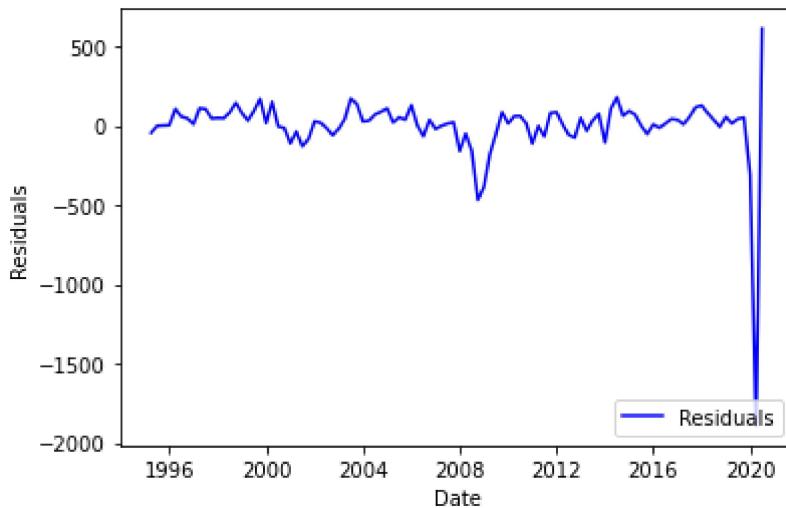
```
# 1. Check model estimates
arma1_1 = tsa.arima_model.ARMA(diff_data,order=(1,1)).fit()
print(arma1_1.summary())
```

```
ARMA Model Results
=====
Dep. Variable: Real GDP No. Observations: 1
02
Model: ARMA(1, 1) Log Likelihood -695.4
34
Method: css-mle S.D. of innovations 221.0
65
Date: Sun, 08 Nov 2020 AIC 1398.8
68
Time: 11:23:17 BIC 1409.3
68
Sample: 04-01-1995 HQIC 1403.1
19
- 07-01-2020
=====
coef std err z P>|z| [0.025
0.975]
-----
const 75.9431 16.595 4.576 0.000 43.418 1
08.468
ar.L1.Real GDP -0.3505 0.574 -0.610 0.542 -1.476
0.775
ma.L1.Real GDP 0.0171 0.530 0.032 0.974 -1.022
1.057
Roots
=====
Real Imaginary Modulus Frequenc
y
-----
AR.1 -2.8532 +0.0000j 2.8532 0.500
0
MA.1 -58.6001 +0.0000j 58.6001 0.500
0
-----
```

C:\Users\35387\anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:159: ValueWarning: No frequency information was provided, so inferred frequency QS-OCT will be used.
`warnings.warn('No frequency information was'

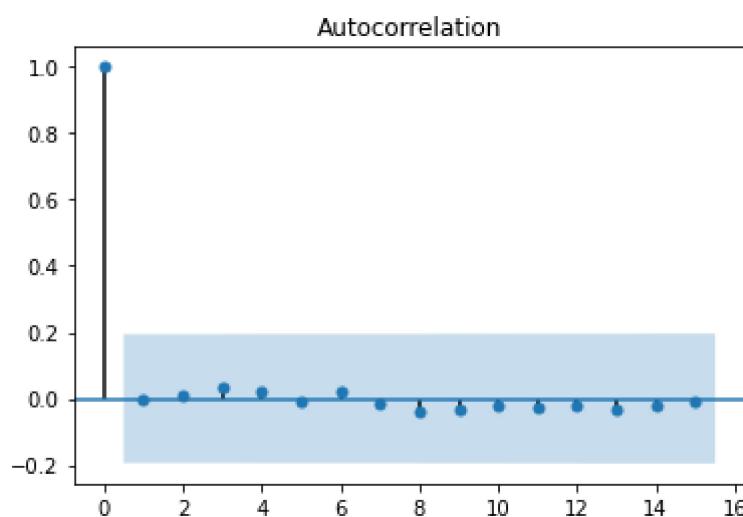
In [192]:

```
# plot residual errors
res1_1 = pd.DataFrame(arma1_1.resid)
plt.plot(res1_1, 'b-', label='Residuals')
# Set the Labels for axies
plt.xlabel('Date')
plt.ylabel('Residuals')
# Set the Legend position
plt.legend(loc='lower right')
plt.show()
```



In [196]:

```
# # Plot the acf of residuals
plot_acf(res1_1, lags=15)
plt.show()
```

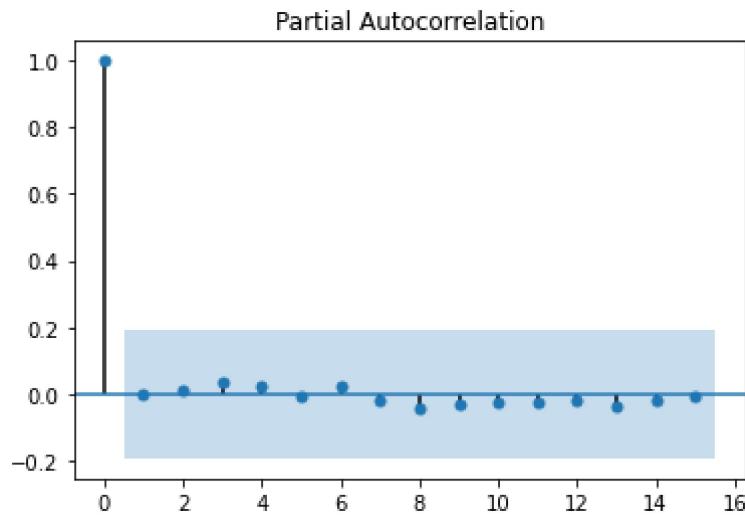


In [164]:

```
#Seems to be no autocorrelation of the residuals
```

In [197]:

```
# # Plot the pacf of residuals
plot_pacf(res1_1, lags=15)
plt.show()
```



In [198]:

```
#Seems to be no autocorrelation of the residuals
#Check non-zero autocorrelation Lags with Ljung-Box test:
```

```
# Check possible non-zero lags: i.e. lag 2 in ARMA(1,1)
lag2_name = ['Ljung-Box Test Statistic', 'p-value']
lag2_lb = acorr_ljungbox(res1_1, lags=2)
# Indexing the stats and p-value for the 2th lag
lag2_q = lag2_lb[0][-1]
lag2_p = lag2_lb[1][-1]
# Q-stats
q_stats_lag2 = {'Ljung-Box Test Statistic':lag2_q, 'p-value':lag2_p}
q_stats_lag2
```

```
C:\Users\35387\anaconda3\lib\site-packages\statsmodels\stats\diagnostic.py:5
24: FutureWarning: The value returned will change to a single DataFrame after 0.12 is released. Set return_df to True to use to return a DataFrame now.
Set return_df to False to silence this warning.
warnings.warn(msg, FutureWarning)
```

Out[198]:

```
{'Ljung-Box Test Statistic': 0.014813203091533458,
 'p-value': 0.9926207597344222}
```

In [176]:

```
#ARMA(2,1)
```

In [199]:

```
# 1. Check model estimates
arma2_1 = tsa.arima_model.ARMA(diff_data,order=(2,1)).fit()
print(arma2_1.summary())
```

ARMA Model Results

```
=====
===
Dep. Variable:          Real GDP    No. Observations:                 1
02
Model:                  ARMA(2, 1)    Log Likelihood:                -695.3
49
Method:                 css-mle     S.D. of innovations:           220.8
41
Date:                   Sun, 08 Nov 2020   AIC:                         1400.6
99
Time:                   13:16:53      BIC:                         1413.8
24
Sample:                 04-01-1995   HQIC:                        1406.0
13
                           - 07-01-2020
=====
```

```
=====
=====
```

	coef	std err	z	P> z	[0.025
0.975]					
const	74.4998	18.880	3.946	0.000	37.496
11.504					
ar.L1.Real GDP	0.2266	0.674	0.336	0.737	-1.094
1.548					
ar.L2.Real GDP	0.2686	0.252	1.066	0.287	-0.226
0.763					
ma.L1.Real GDP	-0.5741	0.652	-0.880	0.379	-1.852
0.704					

Roots

```
=====
=====
```

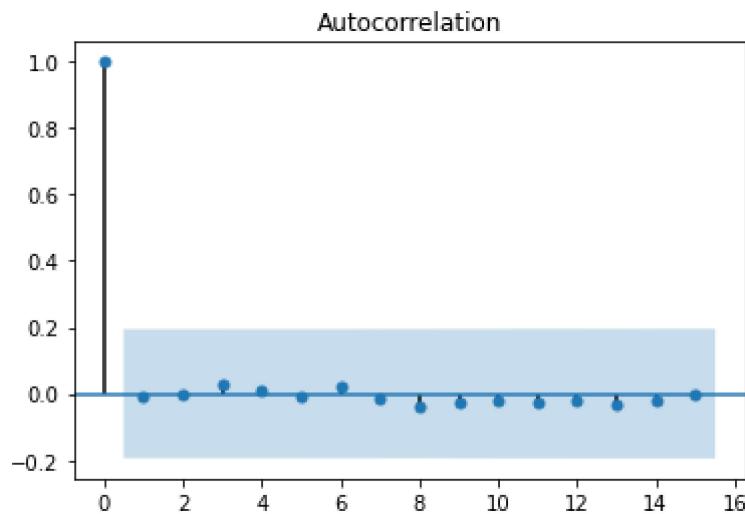
	Real	Imaginary	Modulus	Frequenc
y				
-				
AR.1	1.5532	+0.0000j	1.5532	0.000
0				
AR.2	-2.3966	+0.0000j	2.3966	0.500
0				
MA.1	1.7418	+0.0000j	1.7418	0.000
0				

```
C:\Users\35387\anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.p
y:159: ValueWarning: No frequency information was provided, so inferred freq
uency QS-OCT will be used.
```

```
warnings.warn('No frequency information was'
```

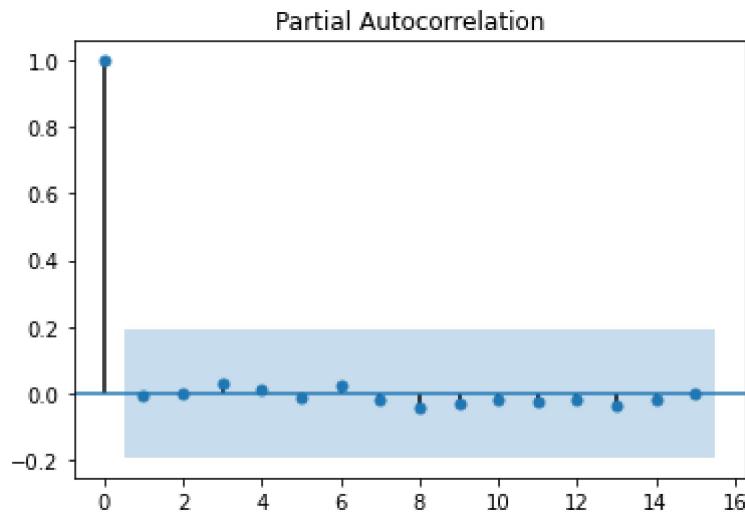
In [200]:

```
# # 2. Check correlogram of residuals: ARMA(2,1)
res2_1 = pd.DataFrame(arma2_1.resid)
plot_acf(res2_1, lags=15)
plt.show()
```



In [201]:

```
plot_pacf(res2_1, lags=15)
plt.show()
```



In [202]:

```
#Seems to be no autocorrelation of the residuals  
#Check non-zero autocorrelation Lags with with Ljung-Box test:
```

```
# Check possible non-zero Lags: i.e. Lag 8 in ARMA(2,1)  
lag8_name = ['Ljung-Box Test Statistic', 'p-value']  
lag8_lb = acorr_ljungbox(res2_1, lags=8)  
# Indexing the stats and p-value for the 8th lag  
lag8_q = lag8_lb[0][-1]  
lag8_p = lag8_lb[1][-1]  
# Q-stats  
q_stats_lag8 = {'Ljung-Box Test Statistic':lag8_q, 'p-value':lag8_p}  
q_stats_lag8
```

```
C:\Users\35387\anaconda3\lib\site-packages\statsmodels\stats\diagnostic.py:5  
24: FutureWarning: The value returned will change to a single DataFrame afte  
r 0.12 is released. Set return_df to True to use to return a DataFrame now.  
Set return_df to False to silence this warning.  
    warnings.warn(msg, FutureWarning)
```

Out[202]:

```
{'Ljung-Box Test Statistic': 0.3432142272253769, 'p-value': 0.99996848772963  
24}
```

In [203]:

1. Check model estimates

```
arma7_1 = tsa.arima_model.ARMA(diff_data,order=(7,1)).fit()
print(arma7_1.summary())
```

C:\Users\35387\anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.p
y:159: ValueWarning: No frequency information was provided, so inferred freq
uency QS-OCT will be used.

```
warnings.warn('No frequency information was'
```

ARMA Model Results

```
=====
Dep. Variable: Real GDP No. Observations: 1
02
Model: ARMA(7, 1) Log Likelihood -693.5
21
Method: css-mle S.D. of innovations 214.1
04
Date: Sun, 08 Nov 2020 AIC 1407.0
43
Time: 13:26:26 BIC 1433.2
93
Sample: 04-01-1995 HQIC 1417.6
72
- 07-01-2020
=====
```

	coef	std err	z	P> z	[0.025
0.975]					

	coef	std err	z	P> z	[0.025
0.975]					
const	75.2593	5.469	13.762	0.000	64.541
85.978					
ar.L1.Real GDP	0.6320	0.120	5.284	0.000	0.398
0.866					
ar.L2.Real GDP	0.2625	0.312	0.842	0.400	-0.349
0.874					
ar.L3.Real GDP	0.1951	0.324	0.603	0.547	-0.439
0.829					
ar.L4.Real GDP	-0.0209	0.336	-0.062	0.950	-0.680
0.638					
ar.L5.Real GDP	-0.2250	0.331	-0.679	0.497	-0.874
0.424					
ar.L6.Real GDP	0.2165	0.305	0.709	0.478	-0.382
0.815					
ar.L7.Real GDP	-0.1816	0.220	-0.827	0.409	-0.612
0.249					
ma.L1.Real GDP	-1.0000	0.036	-27.736	0.000	-1.071
-0.929					

Roots

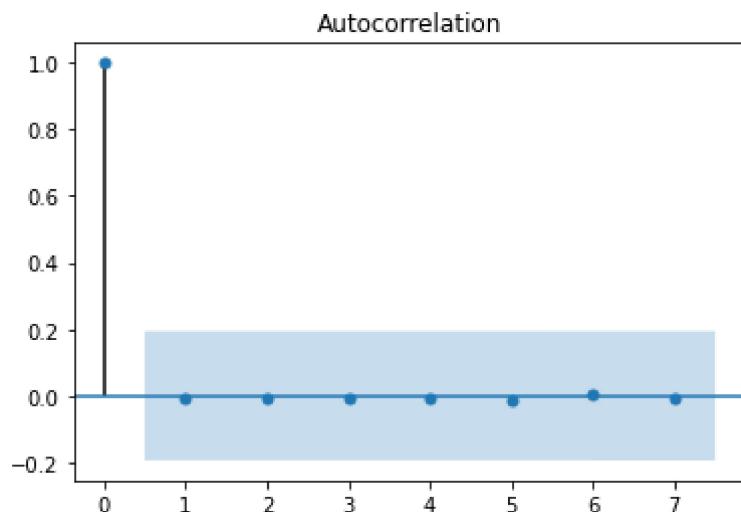
```
=====
Real Imaginary Modulus Frequenc
y
-----
- AR.1 -1.1882 -0.0000j 1.1882 -0.500
0
AR.2 -0.5734 -1.0578j 1.2032 -0.329
```

1				
AR.3	-0.5734	+1.0578j	1.2032	0.329
1				
AR.4	1.1204	-0.1544j	1.1310	-0.021
8				
AR.5	1.1204	+0.1544j	1.1310	0.021
8				
AR.6	0.6431	-1.4454j	1.5820	-0.183
4				
AR.7	0.6431	+1.4454j	1.5820	0.183
4				
MA.1	1.0000	+0.0000j	1.0000	0.000
0				

-
◀ ▶

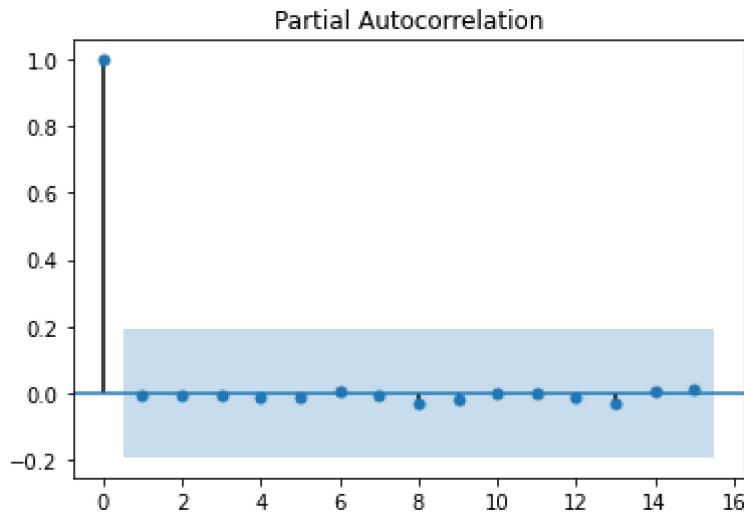
In [204]:

```
# 2. Check correlogram of residuals: ARMA(7,1)
res7_1 = pd.DataFrame(arma7_1.resid)
plot_acf(res7_1, lags=7)
plt.show()
```



In [205]:

```
plot_pacf(res7_1, lags=15)
plt.show()
```



In [206]:

```
# Summarizing all IC
IC = {'AIC':[arma1_1.aic,arma2_1.aic,arma7_1.aic], 'BIC':[arma1_1.bic,arma2_1.bic,arma7_1.bi
    'HQIC':[arma1_1.hqic,arma2_1.hqic,arma7_1.hqic]}
IC_df = pd.DataFrame(IC,index=['ARMA(1,1)', 'ARMA(2,1)', 'ARMA(7,1)'])
print(IC_df)
```

	AIC	BIC	HQIC
ARMA(1,1)	1398.867629	1409.367520	1403.119393
ARMA(2,1)	1400.698783	1413.823647	1406.013488
ARMA(7,1)	1407.042999	1433.292727	1417.672409